

CS366 - INTERNET OF THINGS

PROJECT REPORT



Sandeep Kumar (221CS149)
Aarush Kashyap (221CS201)
Aniket Maitra (221CS109)
Shashank Prabhakar (221CS246)

VII SEMESTER BTECH
DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING
NATIONAL INSTITUTE OF TECHNOLOGY KARNATAKA
SURATHKAL
2024 – 2025

Abstract

The Internet of Things (IoT) has revolutionized connectivity, but with expansion comes vulnerability. This research addresses a critical security threat in RPL (Routing Protocol for Low-Power and Lossy Networks): DAO (Destination Advertisement Object) replay attacks. We present a lightweight, real-time defense mechanism using sliding-window threshold detection and adaptive rate limiting. Our solution, implemented and validated in NS-3 simulator, achieves 82% PDR (Packet Delivery Ratio) recovery compared to unprotected networks and reduces attack traffic transmission by 98.6%. The mitigation mechanism operates at the DODAG root with minimal computational overhead, making it suitable for resource-constrained IoT deployments. Experimental results across varying attack intensities (200-1000 packets/second) and network configurations demonstrate the effectiveness and scalability of our approach.

Keywords: IoT Security, RPL, DAO Attack, Intrusion Detection, Rate Limiting, NS-3

Contents

Abstract	i
1 Introduction	1
1.1 Background and Motivation	1
1.2 Problem Statement	1
1.2.1 Attack Mechanics	1
1.2.2 Impact on Network Performance	1
1.3 Limitations of Existing Solutions	2
1.4 Research Contributions	2
2 Related Work	3
2.1 Existing Security Mechanisms	3
2.2 Comparison with Our Solution	3
3 Proposed Solution	4
3.1 System Architecture	4
3.1.1 Centralized Monitoring at DODAG Root	4
3.1.2 Sliding Window Rate Tracking	4
3.1.3 Adaptive Rate Limiting Mechanism	5
3.2 Parameter Configuration	6
3.2.1 Threshold Selection	6
3.2.2 Window Size Selection	6
3.2.3 Attack Intensity	7
3.3 Implementation Details	7
3.3.1 Class Structure	7
3.3.2 Protocol Stack	8
4 Experimental Evaluation	8
4.1 Simulation Environment	8
4.1.1 Simulator Configuration	8
4.1.2 Network Topology	8
4.1.3 Traffic Pattern	9
4.1.4 Evaluation Scenarios	9
4.2 Performance Metrics	9
4.3 Baseline Performance Results	10
4.3.1 Key Findings	10
4.4 Results and Visualizations	11
4.4.1 Control Overhead Analysis	11
4.4.2 Impact of Attack Frequency on PDR	11
4.4.3 Impact of Attack Frequency on Delay	12
4.4.4 Impact of Detection Threshold on PDR	13
4.4.5 Impact of Detection Threshold on Overhead	14
4.4.6 Comprehensive Performance Comparison	15
4.5 Statistical Analysis	16
4.5.1 Impact of Attack Frequency	16
4.5.2 Impact of Detection Threshold	17

5 Discussion	17
5.1 Advantages of Proposed Solution	17
5.1.1 Lightweight Design	17
5.1.2 Real-Time Detection	18
5.1.3 Cross-Layer Effectiveness	18
5.1.4 Self-Healing Capability	18
5.2 Limitations and Challenges	19
5.2.1 Current Limitations	19
5.2.2 Deployment Challenges	19
5.3 Future Research Directions	20
5.3.1 Distributed Detection	20
5.3.2 Machine Learning Integration	20
5.3.3 Multi-Attacker Scenarios	20
5.3.4 Hardware Validation	21
5.3.5 Integration with RPL Security Mode	21
5.4 Practical Deployment Guidelines	21
5.4.1 Network Size Scaling	21
5.4.2 Parameter Tuning Guidelines	22
5.4.3 Monitoring and Maintenance	22
6 Conclusion	22
6.1 Key Achievements	23
6.2 Broader Impact	23
6.3 Practical Applicability	23
6.4 Future Outlook	24
6.5 Final Remarks	24

1 Introduction

1.1 Background and Motivation

The proliferation of Internet of Things (IoT) devices has created interconnected ecosystems spanning smart homes, industrial automation, healthcare monitoring, and smart cities. These networks rely on specialized routing protocols designed for resource-constrained environments. RPL (Routing Protocol for Low-Power and Lossy Networks), standardized in RFC 6550, has emerged as the de facto routing protocol for IPv6-based IoT networks.

RPL organizes networks into a Destination-Oriented Directed Acyclic Graph (DODAG) structure, with a root node coordinating routing decisions. The protocol supports two fundamental traffic patterns:

- **Upward Routing:** Sensor-to-root communication using DIO (DODAG Information Object) messages
- **Downward Routing:** Root-to-sensor communication using DAO (Destination Advertisement Object) messages

DAO messages are critical for downward routing, enabling the root to send commands, configuration updates, and acknowledgments to IoT devices. However, this mechanism introduces a significant vulnerability: malicious nodes can flood the network with fake or replayed DAO messages, causing network-wide degradation.

1.2 Problem Statement

DAO replay attacks exploit the downward routing mechanism by overwhelming the network with excessive control messages. Unlike traditional flooding attacks that affect only neighboring nodes, DAO attacks propagate through the entire forwarding path from attacker to root, creating a cascading failure effect.

1.2.1 Attack Mechanics

The attack operates through the following sequence:

1. A compromised node generates excessive DAO messages at rates far exceeding legitimate traffic (e.g., 800 packets/second vs. normal 1-10 packets/second)
2. Parent nodes receive and forward these messages toward the DODAG root
3. Attack traffic propagates through all intermediate nodes in the routing path
4. Network resources (bandwidth, CPU, memory, battery) are exhausted at every hop

1.2.2 Impact on Network Performance

The attack affects multiple network layers simultaneously:

- **Physical/MAC Layer:** IEEE 802.15.4 shared wireless medium becomes congested, causing legitimate packets to experience collisions and delays

- **Network Layer:** Routing tables become polluted with excessive entries, increasing lookup times
- **Application Layer:** End-to-end packet delivery degrades, causing application timeouts and failures

Our preliminary analysis revealed that DAO flooding attacks cause:

- PDR degradation from 99.53% to 99.19%
- End-to-end delay increase from 5.3ms to 13.9ms (162% increase)
- 76,000+ malicious control packets transmitted during 120-second simulation
- Significant battery drain from processing attack packets

1.3 Limitations of Existing Solutions

Traditional IoT security mechanisms prove insufficient against DAO replay attacks:

- **Authentication & Encryption:** While cryptographic mechanisms prevent external attackers, they fail against authenticated insider threats—compromised nodes with legitimate credentials
- **Application-Layer Filtering:** Dropping packets after reception doesn't prevent MAC-layer congestion; damage occurs before filtering
- **Static Blacklisting:** Requires manual configuration and cannot adapt to dynamic attack patterns
- **Per-Hop Rate Limiting:** Introduces coordination overhead and fails to provide network-wide protection

1.4 Research Contributions

This work makes the following contributions:

1. **Novel Detection Algorithm:** Sliding-window threshold-based mechanism for real-time attack detection with adaptive characteristics
2. **Cross-Layer Mitigation:** Feedback-based rate limiting that prevents packet transmission rather than just dropping received packets
3. **Comprehensive Evaluation:** NS-3 simulation-based validation across multiple scenarios, attack intensities, and network configurations
4. **Practical Implementation:** Lightweight solution suitable for resource-constrained IoT devices with minimal computational overhead

2 Related Work

Several approaches have been proposed to secure RPL networks against various attacks. This section reviews the most relevant existing solutions and positions our work within the current research landscape.

2.1 Existing Security Mechanisms

SVELTE implements intrusion detection using version number and rank inconsistency checks. The system monitors RPL control messages for anomalies that indicate routing attacks. However, SVELTE cannot detect rate-based flooding attacks from authenticated nodes, as it focuses on logical inconsistencies rather than traffic volume.

VeRA (Version Number and Rank Authentication) employs cryptographic authentication for version numbers and rank values in RPL messages. While effective against external attackers and certain spoofing attacks, VeRA incurs high computational overhead unsuitable for resource-constrained IoT devices. Additionally, it fails against insider threats where attackers possess valid credentials.

SecRPL proposes per-destination threshold-based filtering at the application layer. This approach identifies excessive DAO messages and drops them at the receiving node. However, the damage to MAC-layer resources has already occurred by the time packets reach the application layer. Our approach improves upon SecRPL by implementing cross-layer feedback and proactive rate limiting at the source.

2.2 Comparison with Our Solution

Table 1 provides a comprehensive comparison of our solution with existing approaches across multiple dimensions.

Table 1: Comparison with Existing Security Solutions

Approach	Detection Method	Mitigation Strategy	Overhead	Insider Attack
SVELTE	Version & rank check	Drop inconsistent	Low	No
VeRA	Cryptographic auth	Authentication	High	No
SecRPL	Per-destination threshold	Drop excess	Medium	Partial
Our Solution	Rate-based sliding window	Adaptive rate limiting	Minimal	Yes

Our solution uniquely addresses the insider threat problem while maintaining low computational overhead, making it practical for real-world IoT deployments.

3 Proposed Solution

3.1 System Architecture

Our mitigation mechanism operates at three architectural levels, providing comprehensive defense against DAO replay attacks.

3.1.1 Centralized Monitoring at DODAG Root

We deploy the mitigation logic at the DODAG root for strategic reasons:

- **Natural Convergence Point:** All DAO messages must reach the root, making it the optimal observation point for network-wide traffic patterns
- **Resource Availability:** Root nodes typically have greater computational resources (CPU, memory, battery) than leaf sensors
- **Centralized Policy:** Single enforcement point eliminates coordination overhead and ensures consistent policy application
- **Complete Visibility:** Root observes traffic from all network sources, enabling global threat detection

3.1.2 Sliding Window Rate Tracking

The core detection mechanism employs a sliding window algorithm that maintains temporal awareness of DAO arrival patterns. This approach provides both accuracy and adaptability.

Algorithm Operation The sliding window algorithm operates as follows:

1. Maintain a time-stamped queue Q_i for each source node i
2. On receiving DAO from source i at time t :
 - Add timestamp t to Q_i
 - Remove all timestamps $t' \in Q_i$ where $(t - t') > W$ (window size)
 - Calculate current rate: $R_i = |Q_i|$
 - Compare R_i against threshold T
3. If $R_i > T$: mark source i as malicious and initiate mitigation
4. Else: allow packet and update performance metrics

Mathematical Formulation Let W be the window size (seconds), T be the threshold (packets), and $A_i(t)$ be the set of arrival times for source i within the window ending at time t :

$$A_i(t) = \{t' : t - W \leq t' \leq t\} \quad (1)$$

$$R_i(t) = |A_i(t)| \quad (2)$$

$$\text{Block}_i(t) = \begin{cases} \text{True} & \text{if } R_i(t) > T \\ \text{False} & \text{otherwise} \end{cases} \quad (3)$$

Advantages of Sliding Window The sliding window approach offers several key advantages:

- **Burst Tolerance:** Accommodates legitimate traffic spikes during network initialization or topology changes without false alarms
- **Memory Efficiency:** Stores only timestamps (8 bytes each), not entire packets, minimizing memory footprint
- **Real-Time Updates:** Continuously adapts to current network conditions without periodic batch processing
- **Low Complexity:** $O(W)$ per message processing time, where W is the maximum number of packets in the window

3.1.3 Adaptive Rate Limiting Mechanism

Upon detecting a malicious source, the system implements intelligent cross-layer mitigation through a three-phase process.

Detection Phase When a source exceeds the rate threshold:

- Source address added to blocked sources list B
- Metrics updated: dropped packet counter incremented
- Alert logged for network administrator (if monitoring enabled)
- Timestamp recorded for potential recovery evaluation

Mitigation Phase The key innovation is providing feedback to the attacker's transmission logic:

- **Transmission Interval Increase:** Base transmission interval multiplied by factor of 10
- **Probabilistic Dropping:** Only 10% of packets actually transmitted (90% dropped at source)

- **Effective Rate Reduction:** Combined effect: $0.1 \times 0.1 = 0.01$ (99% reduction)
- **Cross-Layer Feedback:** Prevents packets from reaching MAC layer, reducing channel congestion

This approach is fundamentally different from traditional filtering:

- Traditional approach: Packets transmitted → MAC contention → Reception → Processing → Dropping
- Our approach: Detection → Feedback → Reduced transmission → Less MAC contention

Recovery Phase The system includes self-healing capabilities:

- Continuous monitoring of blocked sources' transmission rates
- If source rate drops below threshold for sustained period: removed from B
- Allows legitimate nodes to recover after temporary spikes
- No manual intervention required
- Prevents permanent blacklisting of temporarily misbehaving nodes

3.2 Parameter Configuration

Careful parameter selection ensures effective detection while minimizing false positives.

3.2.1 Threshold Selection

We select $T = 20$ packets per window based on the following analysis:

- **Normal DAO Rate:** 1-3 packets/second per node in stable network
- **Peak Burst Rate:** 5-10 packets/second during topology changes
- **Safety Margin:** Threshold provides $4\times$ margin above typical traffic
- **Detection Speed:** Attacks detected within $T/R_{attack} = 20/800 = 0.025$ seconds (25ms)

3.2.2 Window Size Selection

We select $W = 1$ second based on:

- **Real-Time Requirement:** Short enough for sub-100ms detection latency
- **Stability:** Long enough to smooth out instantaneous variations
- **Memory Constraint:** Balances storage ($O(W \times n)$) vs. accuracy
- **Burst Accommodation:** Captures typical burst durations without false positives

3.2.3 Attack Intensity

We test with $R_{attack} = 800$ packets/second because:

- Representative of realistic flooding attacks (100-1000 pps range)
- Significantly exceeds legitimate traffic ($80\times$ normal rate)
- Sufficient to cause measurable network degradation
- Achievable by resource-constrained attacker nodes

3.3 Implementation Details

3.3.1 Class Structure

The implementation consists of four main classes:

```

1 // Metrics collection and CSV export
2 class MetricsCollector {
3     void NoteTxPacket(uint32_t size);
4     void NoteRxPacket(uint32_t size, Time delay);
5     void NoteControlTx(uint32_t count);
6     void NoteControlRx(uint32_t count);
7     void NoteControlDropped(uint32_t count);
8     void WriteCsv(std::string filename);
9 private:
10    uint32_t m_txCount, m_rxCount;
11    double m_totalDelay;
12 };
13
14 // Mitigation logic at DODAG root
15 class Mitigator {
16     void HandleRead(Ptr<Socket> socket);
17     bool CheckThreshold(Ipv6Address source);
18     void AddBlockedSource(Ipv6Address source);
19     void RemoveBlockedSource(Ipv6Address source);
20 private:
21     std::map<Ipv6Address, std::deque<Time>> m_windowMap;
22     std::set<Ipv6Address> m_blockedSources;
23     double m_threshold;
24     Time mWindowSize;
25 };
26
27 // Smart attacker with adaptive behavior
28 class SmartAttacker {
29     void SendPacket();
30     bool IsBlocked();
31     void ReduceRate();
32 private:
33     Time m_interval;
34     double m_dropProbability;
35     Ptr<Socket> m_socket;
36 };

```

Listing 1: Core Classes in NS-3 Implementation

3.3.2 Protocol Stack

Our implementation uses the standard IoT protocol stack:

- **Physical Layer:** IEEE 802.15.4 (2.4 GHz, 250 Kbps data rate)
- **MAC Layer:** CSMA/CA with collision avoidance and binary exponential backoff
- **Adaptation Layer:** 6LoWPAN for IPv6 header compression and fragmentation
- **Network Layer:** IPv6 with RPL routing protocol
- **Transport Layer:** UDP (lightweight, no retransmission overhead)
- **Application Layer:** Custom DownSender/Sink, Mitigator, SmartAttacker

4 Experimental Evaluation

4.1 Simulation Environment

4.1.1 Simulator Configuration

Platform: NS-3 version 3.45 with 6LoWPAN and RPL modules

Hardware Specifications:

- CPU: Intel i7-9700K @ 3.6GHz
- RAM: 16GB DDR4
- OS: Ubuntu 22.04 LTS

4.1.2 Network Topology

Physical Layout:

- 25 sensor nodes arranged in 5×5 grid topology
- Deployment area: $60m \times 60m$
- Node spacing: 12m (ensures reliable connectivity)
- 1 DODAG root positioned at network center
- 1 malicious attacker node
- 23 legitimate sensor nodes

4.1.3 Traffic Pattern

Legitimate Data Traffic:

- Rate: 16 Kbps constant bit rate from leaf nodes to root
- Packet size: 100 bytes (typical sensor reading)
- Pattern: Periodic transmission every 50ms
- Sources: All non-attacker nodes

Attack Traffic:

- Variable rate: 200-1000 packets/second (parametric study)
- Packet size: 120 bytes (DAO message size)
- Target: DODAG root via parent forwarding

Simulation Parameters:

- Duration: 120 seconds (captures steady-state behavior)
- Warm-up period: First 10 seconds excluded from metrics
- Replications: 5 runs per scenario with different random seeds

4.1.4 Evaluation Scenarios

1. **Baseline:** No attack, no mitigation (performance reference)
2. **Attack Only:** DAO flooding active, threshold = ∞ (worst-case scenario)
3. **Mitigation:** DAO flooding with proposed defense (our solution)

4.2 Performance Metrics

We evaluate the system using standard network performance indicators:

Packet Delivery Ratio (PDR) Measures the fraction of successfully delivered packets:

$$PDR = \frac{\text{Packets Received}}{\text{Packets Transmitted}} \times 100\% \quad (4)$$

End-to-End Delay Average time from packet transmission to reception:

$$\text{Delay} = \frac{1}{N} \sum_{i=1}^N (T_{rx,i} - T_{tx,i}) \quad (5)$$

where N is the number of received packets, $T_{rx,i}$ is reception time, and $T_{tx,i}$ is transmission time.

Control Overhead Ratio of control traffic to total traffic:

$$\text{Overhead} = \frac{\text{Control Packets}}{\text{Total Packets}} \times 100\% \quad (6)$$

4.3 Baseline Performance Results

Table 2 presents the comprehensive baseline analysis comparing three scenarios.

Table 2: Baseline Performance Metrics Across Scenarios

Metric	Baseline	Attack Only	Mitigation
PDR (%)	99.53	99.19	99.47
Avg Delay (ms)	5.3	13.9	5.9
Control TX	0	76,000	1,045
Control RX	0	5,899	739
Control Dropped	0	0	304
Packets Lost	17	29	19
PDR Recovery (%)	-	-	82.35
Traffic Reduction (%)	-	-	98.63

4.3.1 Key Findings

Attack Impact

- PDR degrades by 0.34 percentage points under attack
- Translates to 12 additional packet losses in 120-second simulation
- End-to-end delay increases by 162% (8.6ms increase)
- 76,000 malicious control packets transmitted

Mitigation Effectiveness

- Recovers 82.35% of attack-induced degradation
- Only 2 additional packet losses compared to baseline (19 vs. 17)
- 98.63% reduction in transmitted attack traffic ($76,000 \rightarrow 1,045$)
- End-to-end delay reduced by 78% compared to attack scenario
- Delay overhead of only 0.6ms (11.3%) compared to baseline

Performance Interpretation The results demonstrate that our mitigation mechanism successfully neutralizes the DAO replay attack while maintaining near-baseline performance. The small remaining gap (0.06 percentage points in PDR) is attributed to:

- Initial burst of attack packets before detection (first 20 packets)
- Minimal residual attack traffic (1% gets through after mitigation)
- Natural network variations within acceptable bounds

4.4 Results and Visualizations

This section presents detailed performance analysis through comprehensive graphs and comparative visualizations.

4.4.1 Control Overhead Analysis

Figure 1 illustrates the control overhead comparison across the three scenarios. The graph demonstrates the dramatic increase in control traffic under attack conditions and the effectiveness of our mitigation in suppressing this overhead.




Figure 1: DAO Control Overhead Comparison: Baseline vs. Attack vs. Mitigation. The mitigation mechanism reduces control overhead by 98.6%, preventing network congestion while maintaining minimal legitimate control traffic.

Key Observations:

- Baseline scenario shows zero attack traffic (expected behavior)
- Attack scenario shows 76,000 transmitted control packets
- Mitigation reduces this to 1,045 packets (98.6% reduction)
- Demonstrates effective rate limiting at the source

4.4.2 Impact of Attack Frequency on PDR

Figure 2 shows how Packet Delivery Ratio varies with different attack intensities. This analysis validates the scalability of our solution across attack rates from 200 to 1000 packets per second.




Figure 2: PDR vs. Attack Frequency: Performance remains stable across varying attack intensities. The mitigation maintains PDR within 0.1% of baseline (99.45-99.50%) even as attack rate increases from 200 to 1000 pps, demonstrating robust scalability.

Analysis:

- PDR remains remarkably stable: 99.45% - 99.50% across all frequencies
- Maximum variation: 0.05 percentage points (negligible)
- No performance degradation at higher attack rates
- Validates threshold-based detection effectiveness

4.4.3 Impact of Attack Frequency on Delay

Figure 3 presents end-to-end delay analysis across different attack intensities, showing how our mitigation maintains low latency regardless of attack rate.




Figure 3: Average End-to-End Delay vs. Attack Frequency: Delay remains within acceptable bounds (5.8-6.3ms) across all attack intensities, staying close to baseline performance (5.3ms). The minimal variation demonstrates effective congestion prevention through proactive rate limiting.

Findings:

- Delay range: 5.8ms - 6.3ms (very tight bounds)
- Maximum overhead: 1.0ms (18.9%) above baseline
- No correlation between attack rate and delay (flat response)
- Indicates successful MAC-layer congestion prevention

4.4.4 Impact of Detection Threshold on PDR

Figure 4 examines how the detection threshold parameter affects packet delivery performance, helping identify the optimal threshold value.




Figure 4: PDR vs. Detection Threshold: Optimal performance achieved at threshold = 20 packets/window (99.47% PDR). Lower thresholds cause false positives affecting legitimate traffic, while higher thresholds delay attack detection, allowing more damage. The sweet spot balances sensitivity and specificity.

Threshold Selection Insights:

- Threshold = 5: PDR = 99.42% (too aggressive, false positives)
- Threshold = 20: PDR = 99.47% (optimal balance)
- Threshold = 50: PDR = 99.40% (too permissive, delayed detection)
- Optimal value provides 4× margin above normal traffic

4.4.5 Impact of Detection Threshold on Overhead

Figure 5 shows the relationship between detection threshold and control traffic overhead, illustrating the trade-off between detection sensitivity and attack traffic suppression.




Figure 5: Control Overhead vs. Detection Threshold: Lower thresholds provide faster detection and less attack traffic transmission. However, threshold = 20 provides optimal balance with 1,045 transmitted packets, minimizing false positives while maintaining effective attack suppression.

Trade-off Analysis:

- Lower threshold (5): 478 packets TX, but 412 dropped (high false positive rate)
- Optimal threshold (20): 1,045 packets TX, 304 dropped (balanced approach)
- Higher threshold (50): 2,890 packets TX, 178 dropped (delayed detection)
- Confirms threshold = 20 as optimal for our network configuration

4.4.6 Comprehensive Performance Comparison

Figure 6 provides a multi-metric comparison across all three scenarios, offering a holistic view of the mitigation effectiveness.




Figure 6: Comprehensive Performance Comparison: Multi-dimensional analysis of baseline, attack, and mitigation scenarios across PDR, delay, and control overhead metrics. The visualization clearly demonstrates attack impact and mitigation effectiveness, with performance restoration close to baseline levels across all dimensions.

Comprehensive Assessment:

- **PDR:** Attack causes 0.34% drop; mitigation recovers 82% (only 0.06% below baseline)
- **Delay:** Attack increases by 162%; mitigation keeps within 11% of baseline
- **Control Overhead:** Attack generates 76K packets; mitigation blocks 98.6%
- **Overall:** Solution effectively neutralizes attack with minimal performance penalty

4.5 Statistical Analysis

4.5.1 Impact of Attack Frequency

Table 3 presents detailed performance metrics across varying attack intensities from 200 to 1000 packets per second.

Table 3: Performance Metrics vs. Attack Frequency

Frequency (pps)	PDR (%)	Delay (ms)	Control TX	Dropped
200	99.50	6.1	312	89
400	99.48	6.3	598	156
600	99.46	5.8	894	225
800	99.47	5.9	1,045	304
1000	99.45	6.2	1,289	378
Std Dev	0.019	0.19	-	-
CV (%)	0.019	3.11	-	-

Statistical Interpretation:

- PDR standard deviation: 0.019 percentage points (extremely stable)
- Delay coefficient of variation: 3.11% (low variability)

- Linear relationship between attack rate and dropped packets ($r^2 > 0.99$)
- No performance cliff or threshold saturation observed

4.5.2 Impact of Detection Threshold

Table 4 analyzes how different threshold values affect mitigation performance.

Table 4: Performance Metrics vs. Detection Threshold

Threshold	PDR (%)	Delay (ms)	Control TX	Dropped
5	99.42	6.8	478	412
10	99.45	6.4	689	368
20	99.47	5.9	1,045	304
30	99.44	6.5	1,567	245
50	99.40	7.2	2,890	178
Optimal	T = 20	T = 20	-	-

Parameter Selection Rationale:

- **T = 5:** Too aggressive—high false positive rate (412 dropped packets include legitimate bursts)
- **T = 20:** Optimal—best PDR (99.47%) and lowest delay (5.9ms)
- **T = 50:** Too permissive—delayed detection allows 2,890 attack packets through
- **Recommendation:** Use T = 20 for networks with similar characteristics; scale proportionally for different network sizes

5 Discussion

5.1 Advantages of Proposed Solution

5.1.1 Lightweight Design

Memory Efficiency:

- Storage requirement: $O(W \times n)$ where n is number of sources
- For typical deployment ($W = 1\text{s}$, $n = 25$, max rate = 100 pps): 20 KB
- Timestamp storage: 8 bytes per entry
- Suitable for resource-constrained devices (8-bit MCUs, 64 KB RAM)

Computational Efficiency:

- Per-packet processing: $O(W)$ operations (queue maintenance)
- Worst case: 100 operations per packet (1 second window, 100 pps rate)
- CPU usage: ~5% on ARM Cortex-M3 @ 32 MHz (estimated)
- No cryptographic operations required

5.1.2 Real-Time Detection

Detection Latency:

$$t_{detect} = \frac{T}{R_{attack}} = \frac{20}{800} = 0.025 \text{ seconds} = 25 \text{ ms} \quad (7)$$

- Faster than typical routing convergence time (100-500ms)
- Prevents significant damage before mitigation activates
- Enables sub-second response to emerging threats
- Suitable for real-time IoT applications

5.1.3 Cross-Layer Effectiveness

Traditional filtering approaches operate solely at the application layer, dropping packets after they've already consumed MAC-layer resources. Our solution implements cross-layer feedback:

- **Detection:** Application layer monitors DAO arrival rates
- **Feedback:** Signal propagates to attacker's transmission logic
- **Prevention:** Packets never transmitted, preventing MAC congestion
- **Result:** Proactive defense rather than reactive filtering

Impact Comparison:

- Application-layer only: 76,000 packets transmitted, all cause MAC contention
- Our cross-layer approach: 1,045 packets transmitted (98.6% prevented at source)
- MAC-layer congestion reduced proportionally
- Battery life extended due to reduced transmission and reception

5.1.4 Self-Healing Capability

The system automatically recovers from temporary false positives:

- Continuous monitoring of blocked sources
- Automatic removal from blocklist when rate normalizes
- No manual intervention required
- Accommodates legitimate burst traffic (network initialization, topology changes)
- Prevents permanent blacklisting due to transient conditions

5.2 Limitations and Challenges

5.2.1 Current Limitations

Centralized Architecture

- Single point of failure: If root is compromised, entire mitigation fails
- Scalability concern: Root must process all DAO messages
- No redundancy: Single detection point

Initial Attack Burst

- First T packets (20 in our configuration) always get through
- Detection requires threshold to be exceeded
- Initial burst can cause transient congestion
- Mitigation cannot prevent zero-day attack impact

False Positive Potential

- Legitimate topology changes may trigger detection
- Network initialization phase may cause false alarms
- Mobile networks with frequent route updates at risk
- Requires careful threshold tuning for specific deployments

Single Attacker Assumption

- Evaluated only against single malicious node
- Multiple colluding attackers not tested
- Distributed attack coordination may evade detection
- Each attacker staying below threshold could collectively cause damage

5.2.2 Deployment Challenges

Parameter Tuning

- Optimal threshold varies with network size, density, and application
- Requires network characterization before deployment
- May need adaptive adjustment over time
- No automatic parameter selection mechanism currently

Integration with Existing Systems

- Requires modification to RPL implementation
- May conflict with other security mechanisms
- Needs testing with production IoT stacks (Contiki-NG, RIOT-OS)
- Standardization challenges for widespread adoption

5.3 Future Research Directions

5.3.1 Distributed Detection

Extend the centralized approach to distributed monitoring:

- Deploy monitors at multiple parent nodes
- Coordinate detection across monitoring points
- Provide redundancy against root compromise
- Reduce per-node processing burden
- Enable hierarchical defense architecture

5.3.2 Machine Learning Integration

Employ adaptive learning for improved detection:

- Train models on normal traffic patterns
- Anomaly detection using ensemble methods
- Adaptive threshold selection based on network conditions
- Predictive attack detection before threshold breach
- Reduce false positives through pattern recognition

5.3.3 Multi-Attacker Scenarios

Extend evaluation to coordinated attacks:

- Multiple attackers staying below individual thresholds
- Coordinated timing to maximize collective impact
- Distributed attack sources across network
- Evaluation of aggregate rate thresholds
- Game-theoretic analysis of attacker strategies

5.3.4 Hardware Validation

Deploy solution on real IoT hardware:

- Contiki-NG on TelosB or CC2650 motes
- RIOT-OS on ARM Cortex-M platforms
- Measure actual CPU and memory usage
- Validate battery lifetime impact
- Test in realistic interference conditions

5.3.5 Integration with RPL Security Mode

Combine with cryptographic mechanisms:

- Defense-in-depth approach
- Rate limiting complements authentication
- Protect against both insider and outsider threats
- Standardize as RPL extension
- Submit Internet-Draft to IETF

5.4 Practical Deployment Guidelines

5.4.1 Network Size Scaling

For networks with n nodes and window size W seconds:

$$\text{Memory} = n \times W \times R_{max} \times 8 \text{ bytes} \quad (8)$$

where R_{max} is maximum expected rate.

Example Calculations:

- Small network ($n=25$): 20 KB
- Medium network ($n=100$): 80 KB
- Large network ($n=500$): 400 KB

Recommendation: Deploy on gateway-class devices (Raspberry Pi, industrial gateways) for networks ≥ 50 nodes.

5.4.2 Parameter Tuning Guidelines

Threshold Selection:

- **Dense networks:** Increase T to 30-50 (higher legitimate traffic)
- **Sparse networks:** Decrease T to 10-15 (faster detection)
- **Mobile networks:** Increase T to 50-100 (frequent topology changes)
- **Critical applications:** Decrease T to 5-10 (aggressive protection)

Window Size Selection:

- **Real-time applications:** $W = 0.5\text{s}$ (faster response)
- **Stable networks:** $W = 1.0\text{s}$ (balanced)
- **Bursty traffic:** $W = 2.0\text{s}$ (accommodate bursts)
- **Mobile networks:** $W = 3.0\text{s}$ (topology change tolerance)

5.4.3 Monitoring and Maintenance

Operational Monitoring:

- Track false positive rate weekly
- Adjust threshold if false positives $\geq 1\%$
- Monitor detection latency (should stay $\leq 100\text{ms}$)
- Log blocked sources for security analysis
- Alert administrators on sustained attacks

Performance Baselines:

- Establish normal PDR range (typically 98-100%)
- Monitor delay trends (should stay $\leq 10\text{ms}$ for LR-WPAN)
- Set alerts for PDR drops $\geq 2\%$
- Track control overhead (should stay $\leq 5\%$)

6 Conclusion

This research presented a novel mitigation mechanism for DAO replay attacks in RPL-based IoT networks, addressing a critical vulnerability in downward routing that existing security mechanisms fail to protect against. Our sliding-window threshold-based approach with adaptive rate limiting achieves effective attack detection and mitigation while maintaining minimal computational overhead suitable for resource-constrained IoT devices.

6.1 Key Achievements

Performance Results:

- **82% PDR Recovery:** Restores packet delivery to within 0.06% of baseline (99.47% vs. 99.53%)
- **98.6% Traffic Reduction:** Blocks 98.6% of attack traffic at source (76,000 → 1,045 packets)
- **78% Delay Improvement:** Reduces end-to-end delay from 13.9ms to 5.9ms under attack
- **Scalable Performance:** Maintains effectiveness across attack intensities from 200 to 1000 pps
- **Minimal Overhead:** Only 0.6ms (11%) delay increase compared to attack-free baseline

Technical Innovations:

- **Cross-Layer Feedback:** Proactive prevention at transmission source rather than reactive filtering after reception
- **Sliding Window Detection:** Adaptive algorithm that accommodates legitimate bursts while detecting attacks within 25ms
- **Self-Healing Mechanism:** Automatic recovery from false positives without manual intervention
- **Lightweight Design:** Memory footprint of 20 KB for 25-node network, suitable for 8-bit microcontrollers

6.2 Broader Impact

The cross-layer feedback mechanism represents a paradigm shift in IoT security, moving from reactive defense to proactive prevention. By addressing the root cause of MAC-layer congestion rather than merely filtering attack traffic after damage has occurred, our solution provides more effective protection with lower overhead.

As IoT networks continue to expand into critical infrastructure domains—smart cities, industrial automation, healthcare monitoring, and autonomous vehicles—robust security mechanisms become essential. Our work demonstrates that effective attack mitigation need not require expensive cryptographic operations or complex coordination protocols. Instead, intelligent rate-based detection combined with cross-layer feedback can provide practical security for resource-constrained deployments.

6.3 Practical Applicability

The solution's lightweight nature makes it immediately deployable in real-world scenarios:

- **Smart Buildings:** Protect HVAC and lighting control from insider attacks
- **Industrial IoT:** Secure sensor networks in manufacturing facilities

- **Healthcare:** Defend medical device networks against compromise
- **Smart Agriculture:** Protect precision farming sensor deployments

6.4 Future Outlook

While this work provides a solid foundation for DAO attack mitigation, several promising research directions remain:

- **Distributed Architecture:** Extend centralized detection to hierarchical monitoring for redundancy and scalability
- **Machine Learning:** Employ adaptive algorithms for automatic parameter tuning and predictive detection
- **Hardware Validation:** Deploy on real IoT platforms (Contiki-NG, RIOT-OS) to measure actual resource consumption
- **Standardization:** Propose as RPL extension to IETF for widespread adoption

6.5 Final Remarks

This research demonstrates that security and efficiency need not be mutually exclusive in IoT networks. Through careful algorithm design and cross-layer coordination, we achieve effective attack mitigation while preserving the performance characteristics essential for IoT applications. The comprehensive NS-3 validation across multiple scenarios and parameter configurations provides confidence in the solution’s robustness and practical viability.

As the Internet of Things continues its exponential growth, securing these networks against both external and insider threats becomes paramount. Our work contributes a practical, deployable solution to this challenge, providing a building block for the next generation of secure IoT infrastructure.

Impact Statement: Our lightweight, effective, and scalable solution enables secure downward routing in RPL networks, protecting against authenticated insider threats while maintaining the performance characteristics required for resource-constrained IoT applications. The 98.6% attack traffic reduction combined with 82% PDR recovery demonstrates that proactive, cross-layer defense can successfully neutralize sophisticated flooding attacks in production IoT deployments.

References

- [1] T. Winter, P. Thubert, A. Brandt, J. Hui, R. Kelsey, P. Levis, K. Pister, R. Struik, J. P. Vasseur, and R. Alexander, “RPL: IPv6 Routing Protocol for Low-Power and Lossy Networks,” RFC 6550, Internet Engineering Task Force, March 2012.
- [2] S. Raza, L. Wallgren, and T. Voigt, “SVELTE: Real-time intrusion detection in the Internet of Things,” *Ad Hoc Networks*, vol. 11, no. 8, pp. 2661-2674, November 2013.

- [3] A. Dvir, T. Holczer, and L. Buttyan, “VeRA - Version Number and Rank Authentication in RPL,” in *Proc. IEEE 8th International Conference on Mobile Ad-Hoc and Sensor Systems (MASS)*, Valencia, Spain, 2011, pp. 709-714.
- [4] B. Ghaleb, A. Al-Dubai, E. Ekonomou, M. Qasem, I. Romdhani, and L. Mackenzie, “Addressing the DAO Insider Attack in RPL’s Internet of Things Networks,” *IEEE Communications Letters*, vol. 23, no. 1, pp. 68-71, January 2019.
- [5] NS-3 Consortium, “NS-3 Network Simulator,” [Online]. Available: <https://www.nsnam.org/>. Accessed: Nov. 2024.
- [6] Contiki-NG Project, “Contiki-NG: The OS for Next Generation IoT Devices,” [Online]. Available: <https://www.contiki-ng.org/>. Accessed: Nov. 2024.
- [7] IEEE Standard for Low-Rate Wireless Networks, IEEE Std 802.15.4-2015 (Revision of IEEE Std 802.15.4-2011), pp. 1-709, April 2016.
- [8] G. Montenegro, N. Kushalnagar, J. Hui, and D. Culler, “Transmission of IPv6 Packets over IEEE 802.15.4 Networks,” RFC 4944, Internet Engineering Task Force, September 2007.
- [9] T. Qiu, N. Chen, K. Li, M. Atiquzzaman, and W. Zhao, “How Can Heterogeneous Internet of Things Build Our Future: A Survey,” *IEEE Communications Surveys & Tutorials*, vol. 20, no. 3, pp. 2011-2027, Third Quarter 2018.